

12.

Implement Linear Regression problem. For example, based on a dataset comprising of existing set of prices and area/size of the houses, predict the estimated price of a given house.

```
In [1]: #Required imports
import pandas as pd
import numpy as np
from sklearn import linear_model
import matplotlib.pyplot as plt
```


```
In [2]: # Reading csv file to dataframe
df = pd.read_csv('houseprices.csv')
df.head()
```

```
Out[2]:
```

	area	price
0	2600	550000
1	2800	556000
2	3000	565000
3	3200	610000
4	3400	640000

```
In [3]: # Scatter plot for the dataset
%matplotlib inline
plt.xlabel('area')
plt.ylabel('price')
plt.scatter(df.area,df.price,color='red',marker='+')
```

```
Out[3]: <matplotlib.collections.PathCollection at 0x1c0d22ae3d0>
```



Preparing data for training

```
In [4]: x_df = df.drop('price',axis='columns')
x_df.head()
```

```
Out[4]:
```

	area
0	2600
1	2800
2	3000
3	3200
4	3400

```
In [5]: price = df.price
price
```

```
Out[5]:
```

0	550000
1	556000
2	565000
3	610000
4	640000
5	680000
6	720000
7	725000
8	760000
9	779000
10	800000

Name: price, dtype: int64

Applying Linear Regression

```
In [6]: # Create linear regression object
reg = linear_model.LinearRegression()
reg.fit(x_df,price)
```

```
Out[6]: LinearRegression()
```

```
In [7]: m = reg.coef_
c = reg.intercept_
print('Coefficient, m = ', m)
print('Intercept, c = ', c)

Coefficient, m = [114.12402428]
Intercept, c = 250142.23764093674
```

Predictions

```
In [8]: ans1 = reg.predict([[3300]])
print('(1) Price of a house with area = 3300 sqr ft: ', ans1)

(1) Price of a house with area = 3300 sqr ft: [626751.51777971]
```

```
In [9]: y = m*3300 + c
print('y = m*x + c =', y)
```

```
y = m*x + c = [626751.51777971]
```

Here, we can see that `y = ans1 = 626751.51777971`

Another prediction

```
In [10]: ans2 = reg.predict([[6000]])  
print('(2) Price of a house with area = 6000 sqr ft: ', ans2)
```


(2) Price of a house with area = 6000 sqr ft: [934886.38334779]

```
In [11]: y = m*6000 + c  
print('y = m*x + c =', y)
```

```
y = m*x + c = [934886.38334779]
```

Visualising Best Fit Line

```
In [12]: x = np.linspace(2000,5500)  
y = m*x+c  
plt.plot(x, y, '-r', label='Best Fit Line')  
plt.legend(loc='upper left')  
plt.title('Graph of Best fit line')  
plt.xlabel('area')  
plt.ylabel('price')  
plt.scatter(df.area,df.price,color='red',marker='+')  
plt.grid()  
plt.show()
```



13.

Based on multiple features/variables perform Linear Regression. For example, based on a number of additional features like number of bedrooms, servant room, number of balconies,

number of houses of years a house has been built
– predict the price of a house.

```
In [13]: #Required imports
import pandas as pd
import numpy as np
from sklearn import linear_model
```

```
In [14]: # Reading csv file to dataframe
df = pd.read_csv('houseprices2.csv')
df
```

```
Out[14]:
```

	area	bedrooms	age	price
0	2600	3.0	20	550000
1	3000	4.0	15	565000
2	3200	NaN	18	610000
3	3600	3.0	30	595000
4	4000	5.0	8	760000
5	4100	6.0	8	810000

Data Preprocessing: Fill NA values with median value of a column

```
In [15]: df.bedrooms.median()
```

```
Out[15]: 4.0
```

```
In [16]: df.bedrooms = df.bedrooms.fillna(df.bedrooms.median())
df
```

```
Out[16]:
```

	area	bedrooms	age	price
0	2600	3.0	20	550000
1	3000	4.0	15	565000
2	3200	4.0	18	610000
3	3600	3.0	30	595000
4	4000	5.0	8	760000
5	4100	6.0	8	810000

Applying Linear Regression

```
In [17]: reg = linear_model.LinearRegression()
reg.fit(df.drop('price',axis='columns'), df.price)
```

```
Out[17]: LinearRegression()
```

```
In [18]: m1, m2, m3 = reg.coef_
c = reg.intercept_
print('Coefficients, \
\n\tm1 = {}, \
\n\tm2 = {}, \
\n\tm3 = {}'.format(m1, m2, m3))
print('Intercept, c = ', c)
```

```
Coefficients,
      m1 = 112.06244194213456,
      m2 = 23388.880077939153,
      m3 = -3231.717908632967
Intercept, c = 221323.00186540443
```

Predictions

```
In [19]: ans1 = reg.predict([[3000, 3, 40]])
print('(1) Price of home with 3000 sqr ft area, 3 bedrooms, 40 year old: ', ar
```

```
(1) Price of home with 3000 sqr ft area, 3 bedrooms, 40 year old: [498408.251
58031]
```

```
In [20]: y1 = m1*3000 + m2*3 + m3*40 + c
print('\ty1 = m1*x1 + m2*x2 + m3*x3 + c =\n\t', y1)
```

```
y1 = m1*x1 + m2*x2 + m3*x3 + c =
498408.2515803069
```

```
In [21]: reg.predict([[2500, 5, 10]])
print('(2) Price of home with 2500 sqr ft area, 5 bedrooms, 10 year old: ', ar
```

```
(2) Price of home with 2500 sqr ft area, 5 bedrooms, 10 year old: [934886.383
34779]
```

```
In [22]: y1 = m1*2500 + m2*5 + m3*10 + c
print('\ty1 = m1*x1 + m2*x2 + m3*x3 + c =\n\t', y1)
```

```
y1 = m1*x1 + m2*x2 + m3*x3 + c =
586106.3280241069
```

14.

Implement a classification/ logistic regression problem. For example based on different features of students data, classify, whether a student is suitable for a particular activity. Based on the available dataset, a student can also implement another classification problem like checking whether an email is spam or not.




```
In [23]: # Import and load digits dataset
from sklearn.datasets import load_digits
digits = load_digits()


# Import matplotlib
```

```
%matplotlib inline
import matplotlib.pyplot as plt
```

```
In [24]: # Plot 2D matrix data of digits
plt.gray()
for i in range(5):
    plt.matshow(digits.images[i])
```

<Figure size 432x288 with 0 Axes>





```
In [25]: # Get the attributes/columns of digits dataset
dir(digits)
```

```
Out[25]: ['DESCR', 'data', 'feature_names', 'frame', 'images', 'target', 'target_names']
```

Creating and training the logistic regression model

```
In [26]: # Import the model
from sklearn.linear_model import LogisticRegression
model = LogisticRegression(max_iter=3000)
```

```
In [27]: # Import train_test_split
from sklearn.model_selection import train_test_split

# Split the dataset into training and testing datasets
X_train, X_test, y_train, y_test = train_test_split(digits.data, digits.target,
```

```
In [28]: print(len(X_train), len(X_test), len(y_train), len(y_test))

1437 360 1437 360
```

```
In [29]: # Training the model
model.fit(X_train, y_train)
```

```
Out[29]: LogisticRegression(max_iter=3000)
```

Measuring accuracy of our model

```
In [30]: model.score(X_test, y_test)
```

```
Out[30]: 0.9611111111111111
```

Predictions

```
In [31]: model.predict(digits.data[0:5])
```

```
Out[31]: array([0, 1, 2, 3, 4])
```

```
In [32]: y_predicted = model.predict(X_test)
y_predicted
```

```
Out[32]: array([0, 1, 1, 4, 3, 4, 6, 8, 6, 2, 1, 6, 8, 6, 7, 2, 7, 2, 1, 3, 5, 4,
                3, 9, 4, 9, 6, 7, 7, 7, 1, 6, 8, 5, 6, 4, 4, 1, 4, 4, 4, 2, 9, 3,
                6, 9, 1, 7, 8, 3, 6, 9, 8, 0, 9, 1, 2, 7, 8, 9, 8, 8, 9, 1, 3, 8,
                8, 5, 7, 6, 3, 6, 7, 4, 6, 7, 6, 6, 8, 8, 0, 5, 4, 9, 7, 9, 0, 1,
                8, 7, 2, 1, 6, 3, 8, 0, 2, 1, 1, 5, 0, 0, 7, 3, 5, 1, 5, 6, 0, 5,
                5, 5, 2, 0, 0, 7, 3, 0, 7, 4, 5, 9, 0, 6, 5, 9, 1, 7, 8, 9, 3, 8,
                4, 3, 9, 0, 0, 0, 9, 0, 4, 0, 7, 5, 3, 0, 7, 1, 1, 9, 3, 0, 5, 5,
                3, 7, 6, 8, 7, 9, 8, 7, 6, 5, 9, 4, 8, 2, 6, 2, 9, 3, 0, 4, 6, 9,
                9, 1, 2, 0, 2, 0, 1, 4, 0, 4, 1, 6, 1, 3, 5, 1, 9, 0, 3, 3, 9, 2,
                0, 1, 2, 5, 1, 4, 9, 9, 2, 7, 2, 6, 0, 9, 0, 4, 4, 3, 7, 4, 5, 0,
                4, 4, 7, 0, 9, 3, 1, 4, 3, 6, 5, 7, 2, 3, 5, 7, 2, 9, 7, 4, 2, 1,
                1, 6, 4, 3, 8, 6, 1, 8, 2, 5, 8, 7, 5, 5, 0, 5, 8, 9, 7, 3, 6, 0,
                8, 1, 3, 5, 7, 0, 8, 7, 9, 1, 3, 9, 9, 5, 3, 9, 2, 1, 7, 9, 5, 6,
                3, 3, 0, 7, 9, 5, 8, 4, 5, 8, 1, 8, 1, 9, 4, 9, 3, 7, 7, 3, 7, 9,
                4, 2, 2, 7, 0, 9, 2, 3, 9, 0, 0, 1, 9, 5, 9, 3, 3, 5, 4, 2, 1, 2,
                2, 4, 1, 3, 8, 4, 7, 8, 7, 9, 3, 0, 4, 4, 4, 5, 9, 1, 3, 2, 6, 6,
                9, 0, 0, 6, 2, 7, 3, 5])
```

Confusion matrix

```
In [33]: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_predicted)
cm
```

```
Out[33]: array([[39,  0,  0,  0,  0,  0,  0,  0,  0,  0],
                [ 0, 33,  0,  0,  0,  0,  1,  0,  1,  1],
                [ 0,  1, 29,  0,  0,  0,  0,  0,  0,  0],
                [ 0,  0,  0, 38,  0,  1,  0,  0,  0,  0],
                [ 0,  1,  0,  0, 36,  0,  0,  1,  0,  0],
                [ 0,  0,  0,  0,  0, 32,  0,  0,  0,  0],
                [ 0,  1,  0,  0,  0,  0, 30,  0,  0,  0],
                [ 0,  0,  0,  1,  0,  0,  0, 38,  0,  2],
                [ 0,  1,  0,  0,  0,  0,  0,  0, 29,  0],
                [ 0,  0,  0,  0,  0,  1,  0,  1,  0, 42]], dtype=int64)
```

```
In [34]: import seaborn as sn
plt.figure(figsize = (10,7))
sn.heatmap(cm, annot=True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```


Out[34]: Text(69.0, 0.5, 'Truth')

